

On cyp2d6_substrate, the measured test-set sampling error spans the top 10 leaderboard entries

Wide band = the test-set sampling error WE measured (95% bootstrap CI, $n=135$, half-width 0.126).

Small bars = each model's leaderboard-reported seed std (mean ± 0.025) — about 5 $\times$ tighter than our band.

So the gaps between ranks sit inside sampling noise. (Seed spread vs test-set resampling; we did NOT recompute a per-model CI.)

